


# Priority Queues and Heaps




# Priority Queue ADT



- ◆ A priority queue stores a collection of items
- ◆ An item is a pair (key, element)
- ◆ Main methods of the Priority Queue ADT
  - **insertItem(k, o)**  
inserts an item with key  $k$  and element  $o$
  - **removeMin()**  
removes the item with the smallest key

- ◆ Additional methods
  - **minKey(k, o)**  
returns, but does not remove, the smallest key of an item
  - **minElement()**  
returns, but does not remove, the element of an item with smallest key
  - **size(), isEmpty()**
- ◆ Applications:
  - Standby flyers
  - Auctions
  - Stock market



# Total Order Relation

- ◆ Keys in a priority queue can be arbitrary objects on which an order is defined
- ◆ Two distinct items in a priority queue can have the same key
- ◆ Mathematical concept of total order relation  $\leq$ 
  - **Reflexive** property:  
 $x \leq x$
  - **Antisymmetric** property:  
 $x \leq y \wedge y \leq x \Rightarrow x = y$
  - **Transitive** property:  
 $x \leq y \wedge y \leq z \Rightarrow x \leq z$

# Comparator ADT



- ◆ A *comparator* encapsulates the action of comparing two objects according to a given total order relation
- ◆ A generic priority queue uses a comparator as a template argument, to define the comparison function ( $<, =, >$ )
- ◆ The comparator is external to the keys being compared. Thus, the same objects can be sorted in different ways by using different comparators.
- ◆ When the priority queue needs to compare two keys, it uses its comparator

# Using Comparators in C++



- ◆ A comparator class overloads the "()" operator with a comparison function.
- ◆ Example: Compare two points in the plane lexicographically.

```
class LexCompare {  
public:  
    int operator()(Point a, Point b) {  
        if (a.x < b.x) return -1  
        else if (a.x > b.x) return +1  
        else if (a.y < b.y) return -1  
        else if (a.y > b.y) return +1  
        else return 0;  
    }  
};
```

- ◆ To use the comparator, define an object of this type, and invoke it using its "()" operator:
- ◆ Example of usage:

```
Point p(2.3, 4.5);  
Point q(1.7, 7.3);  
LexCompare lexCompare;  
  
if (lexCompare(p, q) < 0)  
    cout << "p less than q";  
else if (lexCompare(p, q) == 0)  
    cout << "p equals q";  
else if (lexCompare(p, q) > 0)  
    cout << "p greater than q";
```

# Sorting with a Priority Queue



- ◆ We can use a priority queue to sort a set of comparable elements
  - Insert the elements one by one with a series of **insertItem**(*e*, *e*) operations
  - Remove the elements in sorted order with a series of **removeMin**() operations
- ◆ The running time of this sorting method depends on the priority queue implementation

## Algorithm **PQ-Sort**(*S*, *C*)

**Input** sequence *S*, comparator *C* for the elements of *S*

**Output** sequence *S* sorted in increasing order according to *C*

*P* ← priority queue with comparator *C*

**while** !*S.isEmpty* ()

*e* ← *S.remove* (*S.first* ())

*P.insertItem*(*e*, *e*)

**while** !*P.isEmpty*()

*e* ← *P.minElement*()

*P.removeMin*()

*S.insertLast*(*e*)

# Sequence-based Priority Queue

- ◆ Implementation with an unsorted list



- ◆ Performance:

- **insertItem**
  - ◆ takes  $O(1)$  time since we can insert the item at the beginning or end of the sequence
- **removeMin, minKey and minElement**
  - ◆ take  $O(n)$  time since we have to traverse the entire sequence to find the smallest key


- ◆ Implementation with a sorted list



- ◆ Performance:

- **insertItem**
  - ◆ takes  $O(n)$  time since we have to find the place where to insert the item
- **removeMin, minKey and minElement**
  - ◆ take  $O(1)$  time since the smallest key is at the beginning of the sequence

# Selection-Sort



- ◆ Selection-sort is the variation of PQ-sort where the priority queue is implemented with an unsorted sequence



- ◆ Running time of Selection-sort:
  - Inserting the elements into the priority queue with  $n$  **insertItem** operations takes  $O(n)$  time
  - Removing the elements in sorted order from the priority queue with  $n$  **removeMin** operations takes time proportional to

$$1 + 2 + \dots + n$$

- ◆ Selection-sort runs in  $O(n^2)$  time



# Insertion-Sort



- ◆ Insertion-sort is the variation of PQ-sort where the priority queue is implemented with a sorted sequence



- ◆ Running time of Insertion-sort:
  - Inserting the elements into the priority queue with  $n$  **insertItem** operations takes time proportional to
$$1 + 2 + \dots + n$$
  - Removing the elements in sorted order from the priority queue with a series of  $n$  **removeMin** operations takes  $O(n)$  time
- ◆ Insertion-sort runs in  $O(n^2)$  time

# What is a heap?



- ◆ A heap is a binary tree storing keys at its internal nodes and satisfying the following properties:


- **Heap-Order:**

- ◆ for every internal node  $v$  other than the root,  
 $key(v) \geq key(parent(v))$

- **Complete Binary Tree:**

- ◆ let  $h$  be the height of the heap
- ◆ for  $i = 0, \dots, h - 1$ , there are  $2^i$  nodes of depth  $i$
- ◆ at depth  $h - 1$ , the internal nodes are to the left of the external nodes

- ◆ The last node of a heap is the rightmost internal node of depth  $h - 1$




# Height of a Heap



◆ **Theorem:** A heap storing  $n$  keys has height  $O(\log n)$


Proof: (we apply the complete binary tree property)

- Let  $h$  be the height of a heap storing  $n$  keys
- Since there are  $2^i$  keys at depth  $i = 0, \dots, h-2$  and at least one key at depth  $h-1$ , we have  $n \geq 1 + 2 + 4 + \dots + 2^{h-2} + 1$
- Thus,  $n \geq 2^{h-1}$ , i.e.,  $h \leq \log n + 1$



# Heaps and Priority Queues



- ◆ We can use a heap to implement a priority queue
- ◆ We store a (key, element) item at each internal node
- ◆ We keep track of the position of the last node
- ◆ For simplicity, we show only the keys in the pictures



# Insertion into a Heap




- ◆ Method `insertItem` of the priority queue ADT corresponds to the insertion of a key  $k$  to the heap
- ◆ The insertion algorithm consists of three steps
  - Find the insertion node  $z$  (the new last node)
  - Store  $k$  at  $z$  and expand  $z$  into an internal node
  - Restore the heap-order property (discussed next)





# Upheap

- ◆ After the insertion of a new key  $k$ , the heap-order property may be violated
- ◆ Algorithm upheap restores the heap-order property by swapping  $k$  along an upward path from the insertion node
- ◆ Upheap terminates when the key  $k$  reaches the root or a node whose parent has a key smaller than or equal to  $k$
- ◆ Performance
  - Since a heap has height  $O(\log n)$ , upheap runs in  $O(\log n)$  time





# Removal from a Heap

- ◆ Method `removeMin` of the priority queue ADT corresponds to the removal of the root key from the heap
- ◆ The removal algorithm consists of three steps
  - Replace the root key with the key of the last node  $w$
  - Compress  $w$  and its children into a leaf
  - Restore the heap-order property (discussed next)




# Downheap

- ◆ After replacing the root key with the key  $k$  of the last node, the heap-order property may be violated
- ◆ Algorithm downheap restores the heap-order property by swapping key  $k$  along a downward path from the root
- ◆ Upheap terminates when key  $k$  reaches a leaf or a node whose children have keys greater than or equal to  $k$
- ◆ Performance
  - Since a heap has height  $O(\log n)$ , downheap runs in  $O(\log n)$  time







# Example




# Example




# Example




# Example



# Example



# Example



# Accessing the Queue

- ◆ In a regular queue, you can explicitly keep
  - the head-index and tail-index, or
  - the head-index and the size
- ◆ In a priority queue, you can explicitly keep
  - the head-pointer (root) and the tail-pointer (last node), or
  - the head-pointer and the size


# Question:

- ◆ How do you update the last node (“tail”) pointer or get it from the queue size?
- ◆ Two answers...



# 1=Updating Last Node Pointer

- ◆ The insertion node can be found by traversing a path:
  - Go up until a left child or the root is reached
  - If a left child is reached, go to the right child
  - Go down left until a leaf is reached
- ◆ Similar algorithm for updating the last node after a removal
- ◆ **Performance:**
  - $O(\log n)$




- [illegible]

# Examples



# Heap-Sort



◆ Consider a priority queue with  $n$  items implemented by means of a heap

- the space used is
  - ◆  $O(n)$
- methods **insertItem** and **removeMin** take time
  - ◆  $O(\log n)$
- methods **size**, **isEmpty**, **minKey**, and **minElement** take time
  - ◆  $O(1)$


◆ Using a heap-based priority queue, we can sort a sequence of  $n$  elements in time

- $O(n \log n)$

◆ The resulting algorithm is called heap-sort

◆ Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort

# Heap-Sort




- ◆ More explicitly, how much time does it take to construct a heap?
  - $n$  items, each requiring up to  $\log n$  swaps during “up-heap” operations
    - ◆  $O(n \log n)$
- ◆ How much time does it take to “destruct” a heap (or remove items in sorted order)
  - $n$  items, each requiring up to  $\log n$  swaps during “down-heap” operations
    - ◆  $O(n \log n)$
- ◆ Thus Heap-Sort is
  - $n \log n + n \log n = O(n \log n)$

# Heap Construction

- ◆ Can you do better than  $O(n \log n)$ ?
- ◆ How?
- ◆ Why do we care?
  - We only want to find the few smallest keys among many items
  - We want to quickly start “using the items” in sorted order but the sorting can continue while I start using the first items, e.g.: real-time OS, games, simulations, etc.

# First: Merging Two Heaps


- ◆ We are given two two heaps and a key  $k$
- ◆ We create a new heap with the root node storing  $k$  and with the two heaps as subtrees
- ◆ We perform downheap to restore the heap-order property



# Then: Bottom-up Heap Construction



- ◆ We can construct a heap storing  $n$  given keys using a bottom-up construction with  $\log n$  phases
- ◆ In phase  $i$ , pairs of heaps with  $2^i - 1$  keys are merged into heaps with  $2^{i+1} - 1$  keys








# Example

- ◆ Goal: to create a heap of  $N$  elements
- ◆ Assume  $N = 2^H - 1$  for some integer  $H$  and thus the heap (tree) is “full”
- ◆ In a first step, we construct  $(N+1)/2$  basic heap structures
  - One key and two empty children pointers



# Example





# Example (contd.)




# Example (contd.)




# Example (end)




# Analysis: What is the performance?




- ◆ We visualize the worst-case time of a downheap with a proxy path that goes first right and then repeatedly goes left until the bottom of the heap (this path may differ from the actual downheap path)
- ◆ Since each node is traversed by at most two proxy paths, the total number of nodes of the proxy paths is  $O(n)$ 
  - Or, similarly, each edge of the tree is visited once and since the total number of edges is  $(2n-1)$ , then  $O(n)$
- ◆ Thus, bottom-up heap construction runs in  $O(n)$  time
  - Bottom-up heap construction is faster than  $n$  successive insertions and speeds up the first phase of heap-sort




# Analysis: What is the performance?



- ◆ Thus, we can start using the first results of sorting after  $O(n)$  time and using  $O(n)$  space
  - Groovy!



# Analysis: Why is this important again?




## ◆ Consider the Internet


- You have  $N=10^9$  pages you want to sort and know the top results as soon as possible
- Waiting  $O(N^2) = 10^{18}$  before knowing the top results
  - ◆ is too long...
  - ◆ Even if you choose a very small time unit, e.g.:
    - you assume a 1-GigaHz computer to do 1-Giga operations per second, you will take  $10^9$  seconds, or 31 years!
- Waiting  $O(N \log N) = 30 \times 10^9$ 
  - ◆ is doable, maybe it means 30 seconds
- Waiting  $O(N) = 10^9$ 
  - ◆ **is more doable**, maybe meaning 1 second!!!



# Vector-based Heap Implementation

- ◆ We can represent a heap with  $n$  keys by means of a vector of length  $n + 1$
- ◆ For the node at rank  $i$ 
  - the left child is at rank  $2i$
  - the right child is at rank  $2i + 1$
- ◆ Links between nodes are not explicitly stored
- ◆ The leaves are not represented
- ◆ Operation `insertItem` corresponds to inserting at rank  $n + 1$
- ◆ Operation `removeMin` corresponds to removing at rank  $n$
- ◆ **Yields in-place heap-sort!**





# Lets look at this again

## ◆ Consider the Internet

- We looked at sorting the pages
    - ◆  $O(N^2)$ ,  $O(N \log N)$ ,  $O(N)$  (for first keys of the sort and  $O(N \log N)$  to complete it)
  - How fast can we find any particular page we want in an initially unsorted set?
    - ◆  $O(N^2)$ ?
    - ◆  $O(N \log N)$ ?
    - ◆  $O(N)$ ?
    - ◆  $O(1)$ ?
- ← It is possible! (kinda)

# Coming next!

