


Priority Queues and Heaps




Priority Queue ADT



- ◆ A priority queue stores a collection of items
- ◆ An item is a pair (key, element)
- ◆ Main methods of the Priority Queue ADT
 - **insertItem(k, o)**
inserts an item with key k and element o
 - **removeMin()**
removes the item with the smallest key

- ◆ Additional methods
 - **minKey(k, o)**
returns, but does not remove, the smallest key of an item
 - **minElement()**
returns, but does not remove, the element of an item with smallest key
 - **size(), isEmpty()**
- ◆ Applications:
 - Standby flyers
 - Auctions
 - Stock market



Total Order Relation

- ◆ Keys in a priority queue can be arbitrary objects on which an order is defined
- ◆ Two distinct items in a priority queue can have the same key
- ◆ Mathematical concept of total order relation \leq
 - **Reflexive** property:
 $x \leq x$
 - **Antisymmetric** property:
 $x \leq y \wedge y \leq x \Rightarrow x = y$
 - **Transitive** property:
 $x \leq y \wedge y \leq z \Rightarrow x \leq z$

Comparator ADT



- ◆ A *comparator* encapsulates the action of comparing two objects according to a given total order relation
- ◆ A generic priority queue uses a comparator as a template argument, to define the comparison function ($<, =, >$)
- ◆ The comparator is external to the keys being compared. Thus, the same objects can be sorted in different ways by using different comparators.
- ◆ When the priority queue needs to compare two keys, it uses its comparator

Using Comparators in C++



- ◆ A comparator class overloads the "()" operator with a comparison function.
- ◆ Example: Compare two points in the plane lexicographically.

```
class LexCompare {  
public:  
    int operator()(Point a, Point b) {  
        if (a.x < b.x) return -1  
        else if (a.x > b.x) return +1  
        else if (a.y < b.y) return -1  
        else if (a.y > b.y) return +1  
        else return 0;  
    }  
};
```

- ◆ To use the comparator, define an object of this type, and invoke it using its "()" operator:
- ◆ Example of usage:

```
Point p(2.3, 4.5);  
Point q(1.7, 7.3);  
LexCompare lexCompare;  
  
if (lexCompare(p, q) < 0)  
    cout << "p less than q";  
else if (lexCompare(p, q) == 0)  
    cout << "p equals q";  
else if (lexCompare(p, q) > 0)  
    cout << "p greater than q";
```

Sorting with a Priority Queue



- ◆ We can use a priority queue to sort a set of comparable elements
 - Insert the elements one by one with a series of **insertItem**(e, e) operations
 - Remove the elements in sorted order with a series of **removeMin**() operations
- ◆ The running time of this sorting method depends on the priority queue implementation

Algorithm **PQ-Sort**(*S*, *C*)

Input sequence *S*, comparator *C* for the elements of *S*

Output sequence *S* sorted in increasing order according to *C*

P ← priority queue with comparator *C*

while !*S.isEmpty* ()

e ← *S.remove* (*S.first* ())

P.insertItem(*e*, *e*)

while !*P.isEmpty*()

e ← *P.minElement*()

P.removeMin()

S.insertLast(*e*)

Sequence-based Priority Queue

- ◆ Implementation with an unsorted list



- ◆ Performance:

- **insertItem**
 - ◆ takes $O(1)$ time since we can insert the item at the beginning or end of the sequence
- **removeMin, minKey and minElement**
 - ◆ take $O(n)$ time since we have to traverse the entire sequence to find the smallest key


- ◆ Implementation with a sorted list



- ◆ Performance:

- **insertItem**
 - ◆ takes $O(n)$ time since we have to find the place where to insert the item
- **removeMin, minKey and minElement**
 - ◆ take $O(1)$ time since the smallest key is at the beginning of the sequence

Selection-Sort



- ◆ Selection-sort is the variation of PQ-sort where the priority queue is implemented with an unsorted sequence



- ◆ Running time of Selection-sort:
 - Inserting the elements into the priority queue with n **insertItem** operations takes $O(n)$ time
 - Removing the elements in sorted order from the priority queue with n **removeMin** operations takes time proportional to

$$1 + 2 + \dots + n$$

- ◆ Selection-sort runs in $O(n^2)$ time



Insertion-Sort

- ◆ Insertion-sort is the variation of PQ-sort where the priority queue is implemented with a sorted sequence




- ◆ Running time of Insertion-sort:
 - Inserting the elements into the priority queue with n **insertItem** operations takes time proportional to
$$1 + 2 + \dots + n$$
 - Removing the elements in sorted order from the priority queue with a series of n **removeMin** operations takes $O(n)$ time
- ◆ Insertion-sort runs in $O(n^2)$ time

What is a heap?



- ◆ A heap is a binary tree storing keys at its internal nodes and satisfying the following properties:
 - **Heap-Order:**
 - ◆ for every internal node v other than the root, $key(v) \geq key(parent(v))$
 - **Complete Binary Tree:**
 - ◆ let h be the height of the heap
 - ◆ for $i = 0, \dots, h - 1$, there are 2^i nodes of depth i
 - ◆ at depth $h - 1$, the internal nodes are to the left of the external nodes

- ◆ The last node of a heap is the rightmost internal node of depth $h - 1$




Height of a Heap



◆ **Theorem:** A heap storing n keys has height $O(\log n)$


Proof: (we apply the complete binary tree property)

- Let h be the height of a heap storing n keys
- Since there are 2^i keys at depth $i = 0, \dots, h-2$ and at least one key at depth $h-1$, we have $n \geq 1 + 2 + 4 + \dots + 2^{h-2} + 1$
- Thus, $n \geq 2^{h-1}$, i.e., $h \leq \log n + 1$



Heaps and Priority Queues



- ◆ We can use a heap to implement a priority queue
- ◆ We store a (key, element) item at each internal node
- ◆ We keep track of the position of the last node
- ◆ For simplicity, we show only the keys in the pictures



Insertion into a Heap




- ◆ Method `insertItem` of the priority queue ADT corresponds to the insertion of a key k to the heap
- ◆ The insertion algorithm consists of three steps
 - Find the insertion node z (the new last node)
 - Store k at z and expand z into an internal node
 - Restore the heap-order property (discussed next)





Upheap

- ◆ After the insertion of a new key k , the heap-order property may be violated
- ◆ Algorithm upheap restores the heap-order property by swapping k along an upward path from the insertion node
- ◆ Upheap terminates when the key k reaches the root or a node whose parent has a key smaller than or equal to k
- ◆ Performance
 - Since a heap has height $O(\log n)$, upheap runs in $O(\log n)$ time





Removal from a Heap

- ◆ Method `removeMin` of the priority queue ADT corresponds to the removal of the root key from the heap
- ◆ The removal algorithm consists of three steps
 - Replace the root key with the key of the last node w
 - Compress w and its children into a leaf
 - Restore the heap-order property (discussed next)




Downheap


- ◆ After replacing the root key with the key k of the last node, the heap-order property may be violated
- ◆ Algorithm downheap restores the heap-order property by swapping key k along a downward path from the root
- ◆ Upheap terminates when key k reaches a leaf or a node whose children have keys greater than or equal to k
- ◆ Performance
 - Since a heap has height $O(\log n)$, downheap runs in $O(\log n)$ time




Example




Example




Example




Example



Example



Example



Accessing the Queue


- ◆ In a regular queue, you can explicitly keep
 - the head-index and tail-index, or
 - the head-index and the size
- ◆ In a priority queue, you can explicitly keep
 - the head-pointer (root) and the tail-pointer (last node), or
 - the head-pointer and the size

Question:

- ◆ [TEAMS] How do you update the last node ("tail") pointer or get it from the queue size?
- ◆ Two answers...


1=Updating Last Node Pointer

- ◆ The insertion node can be found by traversing a path:
 - Go up until a left child or the root is reached
 - If a left child is reached, go to the right child
 - Go down left until a leaf is reached
- ◆ Similar algorithm for updating the last node after a removal
- ◆ **Performance:**
 - $O(\log n)$



2=Finding Last Node Pointer


- ◆ The insertion node can be found by traversing a path without needing an explicit tail pointer:
 - Start at the root and using the binary number equivalent of the new number of nodes
 - ◆ Assume the root to be the right-child of an imaginary parent
 - ◆ Starting with MSB, traverse using 0=left and 1=right
 - Prevents the need to keep a last node pointer around
 - Asymptotically same performance, but half the cost
- ◆ Similar algorithm for updating the last node after a removal
- ◆ **Performance:**
 - $O(\log n)$



Examples



Heap-Sort



◆ Consider a priority queue with n items implemented by means of a heap

- the space used is
 - ◆ $O(n)$
- methods **insertItem** and **removeMin** take time
 - ◆ $O(\log n)$
- methods **size**, **isEmpty**, **minKey**, and **minElement** take time
 - ◆ $O(1)$


◆ Using a heap-based priority queue, we can sort a sequence of n elements in time

- $O(n \log n)$

◆ The resulting algorithm is called heap-sort

◆ Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort

Heap-Sort






- ◆ More explicitly, how much time does it take to construct a heap?
 - n items, each requiring up to $\log n$ swaps during “up-heap” operations
 - ◆ $O(n \log n)$
- ◆ How much time does it take to “destruct” a heap (or remove items in sorted order)
 - n items, each requiring up to $\log n$ swaps during “down-heap” operations
 - ◆ $O(n \log n)$
- ◆ Thus Heap-Sort is
 - $n \log n + n \log n = O(n \log n)$

Heap Construction

- ◆ Can you do better than $O(n \log n)$?
- ◆ How?
- ◆ Why do we care?
 - We only want to find the few smallest keys among many items
 - We want to quickly start “using the items” in sorted order but the sorting can continue while I start using the first items, e.g.: real-time OS, games, simulations, etc.

First: Merging Two Heaps


- ◆ We are given two two heaps and a key k
- ◆ We create a new heap with the root node storing k and with the two heaps as subtrees
- ◆ We perform downheap to restore the heap-order property



Then: Bottom-up Heap Construction




- ◆ We can construct a heap storing n given keys using a bottom-up construction with $\log n$ phases
- ◆ In phase i , pairs of heaps with $2^i - 1$ keys are merged into heaps with $2^{i+1} - 1$ keys





Example

- ◆ Goal: to create a heap of N elements
- ◆ Assume $N = 2^H - 1$ for some integer H and thus the heap (tree) is “full”
- ◆ In a first step, we construct $(N+1)/2$ basic heap structures
 - One key and two empty children pointers



Example





Example (contd.)




Example (contd.)




Example (end)




Analysis: What is the performance?




- ◆ We visualize the worst-case time of a downheap with a proxy path that goes first right and then repeatedly goes left until the bottom of the heap (this path may differ from the actual downheap path)
- ◆ Since each node is traversed by at most two proxy paths, the total number of nodes of the proxy paths is $O(n)$
 - Or, similarly, each edge of the tree is visited once and since the total number of edges is $(2n-1)$, then $O(n)$
- ◆ Thus, bottom-up heap construction runs in $O(n)$ time
 - Bottom-up heap construction is faster than n successive insertions and speeds up the first phase of heap-sort




Analysis: What is the performance?



- ◆ Thus, we can start using the first results of sorting after $O(n)$ time and using $O(n)$ space
 - Groovy!



Analysis: Why is this important again?





◆ Consider the Internet

- You have $N=10^9$ pages you want to sort and know the top results as soon as possible
- Waiting $O(N^2) = 10^{18}$ before knowing the top results
 - ◆ is too long...
 - ◆ Even if you choose a very small time unit, e.g.:
 - you assume a 1-GigaHz computer to do 1-Giga operations per second, you will take 10^9 seconds, or 31 years!
- Waiting $O(N \log N) = 30 \times 10^9$
 - ◆ is doable, maybe it means 30 seconds
- Waiting $O(N) = 10^9$
 - ◆ **is more doable**, maybe meaning 1 second!!!

Vector-based Heap Implementation

- ◆ We can represent a heap with n keys by means of a vector of length $n + 1$
- ◆ For the node at rank i
 - the left child is at rank $2i$
 - the right child is at rank $2i + 1$
- ◆ Links between nodes are not explicitly stored
- ◆ The leaves are not represented
- ◆ Operation `insertItem` corresponds to inserting at rank $n + 1$
- ◆ Operation `removeMin` corresponds to removing at rank n
- ◆ **Yields in-place heap-sort!**





Lets look at this again

◆ Consider the Internet

- We looked at sorting the pages
 - ◆ $O(N^2)$, $O(N \log N)$, $O(N)$ (for first keys of the sort and $O(N \log N)$ to complete it)
- How fast can we find any particular page we want in an initially unsorted set?
 - ◆ $O(N^2)$?
 - ◆ $O(N \log N)$?
 - ◆ $O(N)$?
 - ◆ $O(1)$? ← It is possible! (kinda)

Coming next!

